Data Management with SQL for Ecologists (./)

Databases using SQL

Overview

Teaching: 60 min **Exercises:** 5 min

Questions

- What is a relational database and why should I use it?
- · What is SQL?

Objectives

- Describe why relational databases are useful.
- Create and populate a database from a text file.
- · Define SQLite data types.
- Select, group, add to, and analyze subsets of data.
- Combine data across multiple tables.

Setup

Note: this should have been done by participants before the start of the workshop.

We use DB Browser for SQLite (http://sqlitebrowser.org/) and the Portal Project dataset (https://figshare.com/articles/Portal_Project_Teaching_Database/1314459) throughout this lesson. See Setup (../setup.html) for instructions on how to download the data, and also how to install DB Browser for SQLite.

Motivation

To start, let's orient ourselves in our project workflow. Previously, we used Excel and OpenRefine to go from messy, human created data to cleaned, computer-readable data. Now we're going to move to the next piece of the data workflow, using the computer to read in our data, and then use it for analysis and visualization.

What is SQL?

SQL stands for Structured Query Language. SQL allows us to interact with relational databases through queries. These queries can allow you to perform a number of actions such as: insert, update and delete information in a database.

Dataset Description

The data we will be using is a time-series for a small mammal community in southern Arizona. This is part of a project studying the effects of rodents and ants on the plant community that has been running for almost 40 years. The rodents are sampled on a series of 24 plots, with different experimental manipulations controlling which

rodents are allowed to access which plots.

This is a real dataset that has been used in over 100 publications. We've simplified it just a little bit for the workshop, but you can download the full dataset (http://esapubs.org/archive/ecol/E090/118/) and work with it using exactly the same tools we'll learn about today.

Questions

First, let's download and look at some of the cleaned spreadsheets from the Portal Project dataset (https://figshare.com/articles/Portal_Project_Teaching_Database/1314459).

We'll need the following three files:

- surveys.csv
- species.csv
- plots.csv

Challenge

Open each of these csv files and explore them. What information is contained in each file? Specifically, if I had the following research questions:

- How has the hindfoot length and weight of Dipodomys species changed over time?
- What is the average weight of each species, per year?
- What information can I learn about Dipodomys species in the 2000s, over time?

What would I need to answer these questions? Which files have the data I need? What operations would I need to perform if I were doing these analyses by hand?

Goals

In order to answer the questions described above, we'll need to do the following basic data operations:

- select subsets of the data (rows and columns)
- · group subsets of data
- · do math and other calculations
- · combine data across spreadsheets

In addition, we don't want to do this manually! Instead of searching for the right pieces of data ourselves, or clicking between spreadsheets, or manually sorting columns, we want to make the computer do the work.

In particular, we want to use a tool where it's easy to repeat our analysis in case our data changes. We also want to do all this searching without actually modifying our source data.

Putting our data into a relational database and using SQL will help us achieve these goals.

🖈 Definition: Relational Database

A relational database stores data in *relations* made up of *records* with *fields*. The relations are usually represented as *tables*; each record is usually shown as a row, and the fields as columns. In most cases, each record will have a unique identifier, called a *key*, which is stored as one of its fields. Records may also contain keys that refer to records in other tables, which enables us to combine information from two or more sources.

Databases

Why use relational databases

Using a relational database serves several purposes.

- It keeps your data separate from your analysis.
 - This means there's no risk of accidentally changing data when you analyze it.
 - If we get new data we can just rerun the query.
- It's fast, even for large amounts of data.
- It improves quality control of data entry (type constraints and use of forms in MS Access, Filemaker, Oracle Application Express etc.)
- The concepts of relational database querying are core to understanding how to do similar things using programming languages such as R or Python.

Database Management Systems

There are a number of different database management systems for working with relational data. We're going to use SQLite today, but basically everything we teach you will apply to the other database systems as well (e.g. MySQL, PostgreSQL, MS Access, MS SQL Server, Oracle Database and Filemaker Pro). The only things that will differ are the details of exactly how to import and export data and the details of data types.

Relational databases

Let's look at a pre-existing database, the <code>portal_mammals.sqlite</code> file from the Portal Project dataset that we downloaded during Setup (/sql-ecology-lesson/setup/). Click on the "Open Database" button, select the portal mammals.sglite file, and click "Open" to open the database.

You can see the tables in the database by looking at the left hand side of the screen under Database Structure tab. Here you will see a list under "Tables." Each item listed here corresponds to one of the csv files we were exploring earlier. To see the contents of any table, click on it, and then click the "Browse Data" tab next to the "Database Structure" tab. This will give us a view that we're used to - just a copy of the table. Hopefully this helps to show that a database is, in some sense, just a collection of tables, where there's some value in the tables that allows them to be connected to each other (the "related" part of "relational database").

The "Database Structure" tab also provides some metadata about each table. If you click on the down arrow next to a table name, you will see information about the columns, which in databases are referred to as "fields," and their assigned data types.

(The rows of a database table are called *records*.) Each field contains one variety or type of data, often numbers or text. You can see in the surveys table that most fields contain numbers (BIGINT, or big integer, and FLOAT, or floating point numbers/decimals) while the species table is entirely made up of text fields.

The "Execute SQL" tab is blank now - this is where we'll be typing our queries to retrieve information from the database tables.

To summarize:

- Relational databases store data in tables with fields (columns) and records (rows)
- Data in tables has types, and all values in a field have the same type (list of data types)
- Queries let us look up data or make calculations based on columns

Database Design

- Every row-column combination contains a single *atomic* value, i.e., not containing parts we might want to work with separately.
- One field per type of information
- No redundant information
 - Split into separate tables with one table per class of information
 - Needs an identifier in common between tables shared column to reconnect (known as a foreign key).

Import

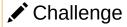
Before we get started with writing our own queries, we'll create our own database. We'll be creating this database from the three csv files we downloaded earlier. Close the currently open database (File > Close Database) and then follow these instructions:

- 1. Start a New Database
 - Click on the New Database icon or select File » New Database
 - Assign a name to the new database, choose the folder where you'd like to save it, and click Save. This
 creates the database in the selected folder.
- 2. Choose Start the import Database -> Import
- 3. We will be importing tables and not creating tables from scratch, so click **Cancel** to edit out of the next pop-up window.
- Select File > Import > Table from CSV file... Choose surveys.csv from the data folder we downloaded and click Open.
- 5. Give the table a name that matches the file name or use the default.
- 6. If the first row has column headings, be sure to check the box next to "Column names in first line."
- 7. Be sure the field separator and quotation options are correct. If you're not sure which options are correct, test some of the options and until the preview at the bottom of the window looks right.
- 8. Click OK
- 9. Back on the Database Structure tab, you should now see the table listed. Right click on the table name and choose **Modify Table**, or click on the **Modify Table** just under the tabs and above the table.
- 10. In the center panel of the windown you'll see, set the data types for each field using the suggestions in the table below (this includes fields from plots and species tables also):

Field	Data Type	Motivation	Table(s)
day	INTEGER	Having data as numeric allows for meaningful arithmetic and comparisons	surveys
genus	TEXT	Field contains text data	species
hindfoot_lengt	hREAL	Field contains measured numeric data	surveys
month	INTEGER	Having data as numeric allows for meaningful arithmetic and comparisons	surveys
plot_id	INTEGER	R Field contains numeric data	plots, surveys
plot_type	TEXT	Field contains text data	plots
record_id	INTEGER	R Field contains numeric data	surveys
sex	TEXT	Field contains text data	surveys
species_id	TEXT	Field contains text data	species,
		riela contains text data	surveys
species	TEXT	Field contains text data	species

Field	Data Type	Motivation	Table(s)
taxa	TEXT	Field contains text data	species
weight	REAL	Field contains measured numerical data	surveys
year	INTEGE	R Allows for meaningful arithmetic and comparisons	surveys

Finally, click **OK** one more time to confirm the operation.



• Import the plots and species tables

You can also use this same approach to append new fields to an existing table.

Adding fields to existing tables

- 1. Go to the "Database Structure" tab, right click on the table you'd like to add data to, and choose **Modify Table**, or click on the **Modify Table** just under the tabs and above the table.
- 2. Click the **Add Field** button to add a new field and assign it a data type.

Data types

Data type	Description
CHARACTER(n)	Character string. Fixed-length n
VARCHAR(n) or CHARACTER VARYING(n)	Character string. Variable length. Maximum length n
BINARY(n)	Binary string. Fixed-length n
BOOLEAN	Stores TRUE or FALSE values
VARBINARY(n) or BINARY VARYING(n)	Binary string. Variable length. Maximum length n
INTEGER(p)	Integer numerical (no decimal).
SMALLINT	Integer numerical (no decimal).
INTEGER	Integer numerical (no decimal).
BIGINT	Integer numerical (no decimal).
DECIMAL(p,s)	Exact numerical, precision p, scale s.
NUMERIC(p,s)	Exact numerical, precision p, scale s. (Same as DECIMAL)
FLOAT(n)	Approximate numerical, mantissa precision p. A floating number in base 10
FLOAT(p)	exponential notation.
REAL	Approximate numerical
FLOAT	Approximate numerical
DOUBLE PRECISION	Approximate numerical
DATE	Stores year, month, and day values
TIME	Stores hour, minute, and second values
TIMESTAMP	Stores year, month, day, hour, minute, and second values
INTERVAL	Composed of a number of integer fields, representing a period of time,
INTERVAL	depending on the type of interval
ARRAY	A set-length and ordered collection of elements
MULTISET	A variable-length and unordered collection of elements
XML	Stores XML data

SQL Data Type Quick Reference

Different databases offer different choices for the data type definition.

The following table shows some of the common names of data types between the various database platforms:

Data type boolean	Access Yes/No	SQLServer Bit	Oracle Byte	MySQL N/A	PostgreSQL Boolean
integer	Number (integer)	Int	Number	Int / Integer	Int / Integer
float	Number (single)	Float / Real	Number	Float	Numeric
currency	Currency	Money	N/A	N/A	Money
string (fixed)	N/A	Char	Char	Char	Char
string (variable)	Text (<256) / Memo (65k+)	Varchar	Varchar	2Varchar	Varchar
binary object OLE Object Memo Binary (fixed up to 8K)	Varbinary (<8K)	Image (<2GB) Long	Raw Blob	Text Binary	Varbinary

Key Points

- SQL allows us to select and group subsets of data, do math and other calculations, and combine data.
- A relational database is made up of tables which are related to each other by shared keys.
- Different database management systems (DBMS) use slightly different vocabulary, but they are all based on the same ideas.

Basic Queries

Overview

Teaching: 30 min **Exercises:** 5 min

Questions

· How do I write a basic query in SQL?

Objectives

- · Write and build queries.
- Filter data given various criteria.
- Sort the results of a query.

Writing my first query

Let's start by using the **surveys** table. Here we have data on every individual that was captured at the site, including when they were captured, what plot they were captured on, their species ID, sex and weight in grams.

Let's write an SQL query that selects only the year column from the surveys table. SQL queries can be written in the box located under the "Execute SQL" tab. Click on the right arrow above the query box to execute the query. (You can also use the keyboard shortcut "Cmd-Enter" on a Mac or "Ctrl-Enter" on a Windows machine to execute a query.) The results are displayed in the box below your query.

```
SELECT year
FROM surveys;
```

We have capitalized the words SELECT and FROM because they are SQL keywords. SQL is case insensitive, but it helps for readability, and is good style.

If we want more information, we can just add a new column to the list of fields, right after SELECT:

```
SELECT year, month, day FROM surveys;
```

Or we can select all of the columns in a table using the wildcard *

```
SELECT *
FROM surveys;
```

Limiting results

Sometimes you don't want to see all the results you just want to get a sense of of what's being returned. In that case you can use the LIMIT command. In particular you would want to do this if you were working with large databases.

```
SELECT *
FROM surveys
LIMIT 10;
```

Unique values

If we want only the unique values so that we can quickly see what species have been sampled we use DISTINCT

```
SELECT DISTINCT species_id
FROM surveys;
```

If we select more than one column, then the distinct pairs of values are returned

```
SELECT DISTINCT year, species_id FROM surveys;
```

Calculated values

We can also do calculations with the values in a query. For example, if we wanted to look at the mass of each individual on different dates, but we needed it in kg instead of g we would use

```
SELECT year, month, day, weight/1000
FROM surveys;
```

When we run the query, the expression weight / 1000 is evaluated for each row and appended to that row, in a new column. If we used the INTEGER data type for the weight field then integer division would have been done, to obtain the correct results in that case divide by 1000.0. Expressions can use any fields, any arithmetic operators

(+, -, *, and /) and a variety of built-in functions. For example, we could round the values to make them easier to read.

```
SELECT plot_id, species_id, sex, weight, ROUND(weight / 1000, 2)
FROM surveys;
```

Challenge

Write a query that returns the year, month, day, species_id and weight in mg.

Solution

```
SELECT day, month, year, species_id, weight * 1000
FROM surveys;
```

Filtering

Databases can also filter data – selecting only the data meeting certain criteria. For example, let's say we only want data for the species *Dipodomys merriami*, which has a species code of DM. We need to add a WHERE clause to our query:

```
SELECT *
FROM surveys
WHERE species_id='DM';
```

We can do the same thing with numbers. Here, we only want the data since 2000:

```
SELECT * FROM surveys
WHERE year >= 2000;
```

If we used the TEXT data type for the year the WHERE clause should be year >= '2000'. We can use more sophisticated conditions by combining tests with AND and OR. For example, suppose we want the data on *Dipodomys merriami* starting in the year 2000:

```
SELECT *
FROM surveys
WHERE (year >= 2000) AND (species_id = 'DM');
```

Note that the parentheses are not needed, but again, they help with readability. They also ensure that the computer combines AND and OR in the way that we intend.

If we wanted to get data for any of the Dipodomys species, which have species codes DM, DO, and DS, we could combine the tests using OR:

```
SELECT *
FROM surveys
WHERE (species_id = 'DM') OR (species_id = 'DO') OR (species_id = 'DS');
```

Challenge

• Produce a table listing the data for all individuals in Plot 1 that weighed more than 75 grams, telling us the date, species id code, and weight (in kg).

```
Solution
```

```
SELECT day, month, year, species_id, weight / 1000
FROM surveys
WHERE (plot_id = 1) AND (weight > 75);
```

Building more complex queries

Now, lets combine the above queries to get data for the 3 *Dipodomys* species from the year 2000 on. This time, let's use IN as one way to make the query easier to understand. It is equivalent to saying

```
WHERE (species_id = 'DM') OR (species_id = 'DO') OR (species_id = 'DS'), but reads more neatly:
```

```
SELECT *
FROM surveys
WHERE (year >= 2000) AND (species_id IN ('DM', 'DO', 'DS'));
```

We started with something simple, then added more clauses one by one, testing their effects as we went along. For complex queries, this is a good strategy, to make sure you are getting what you want. Sometimes it might help to take a subset of the data that you can easily see in a temporary database to practice your queries on before working on a larger or more complicated database.

When the queries become more complex, it can be useful to add comments. In SQL, comments are started by -- , and end at the end of the line. For example, a commented version of the above query can be written as:

```
-- Get post 2000 data on Dipodomys' species
-- These are in the surveys table, and we are interested in all columns
SELECT * FROM surveys
-- Sampling year is in the column `year`, and we want to include 2000
WHERE (year >= 2000)
-- Dipodomys' species have the `species_id` DM, DO, and DS
AND (species_id IN ('DM', 'DO', 'DS'));
```

Although SQL queries often read like plain English, it is *always* useful to add comments; this is especially true of more complex queries.

Sorting

We can also sort the results of our queries by using ORDER BY . For simplicity, let's go back to the **species** table and alphabetize it by taxa.

First, let's look at what's in the **species** table. It's a table of the species_id and the full genus, species and taxa information for each species_id. Having this in a separate table is nice, because we didn't need to include all this information in our main **surveys** table.

```
SELECT *
FROM species;
```

Now let's order it by taxa.

```
SELECT *
FROM species
ORDER BY taxa ASC;
```

The keyword ASC tells us to order it in Ascending order. We could alternately use DESC to get descending order.

```
SELECT *
FROM species
ORDER BY taxa DESC;
```

ASC is the default.

We can also sort on several fields at once. To truly be alphabetical, we might want to order by genus then species.

```
SELECT *
FROM species
ORDER BY genus ASC, species ASC;
```

Challenge

• Write a query that returns year, species_id, and weight in kg from the surveys table, sorted with the largest weights at the top.

Solution

```
SELECT year, species_id, weight / 1000 FROM surveys ORDER BY weight DESC;
```

Order of execution

Another note for ordering. We don't actually have to display a column to sort by it. For example, let's say we want to order the birds by their species ID, but we only want to see genus and species.

```
SELECT genus, species
FROM species
WHERE taxa = 'Bird'
ORDER BY species_id ASC;
```

We can do this because sorting occurs earlier in the computational pipeline than field selection.

The computer is basically doing this:

- 1. Filtering rows according to WHERE
- 2. Sorting results according to ORDER BY
- 3. Displaying requested columns or expressions.

Clauses are written in a fixed order: SELECT, FROM, WHERE, then ORDER BY. It is possible to write a query as a single line, but for readability, we recommend to put each clause on its own line.

- Let's try to combine what we've learned so far in a single query. Using the surveys table write a query to display the three date fields, species_id, and weight in kilograms (rounded to two decimal places), for individuals captured in 1999, ordered alphabetically by the species_id.
- Write the query as a single line, then put each clause on its own line, and see how more legible the query becomes!

Solution

```
SELECT year, month, day, species_id, ROUND(weight / 1000, 2)
FROM surveys
WHERE year = 1999
ORDER BY species_id;
```

• Key Points

- · It is useful to apply conventions when writing SQL queries to aid readability.
- Use logical connectors such as AND or OR to create more complex queries.
- Calculations using mathematical symbols can also be performed on SQL queries.
- Adding comments in SQL helps keep complex queries understandable.

SQL Aggregation and aliases

Overview

Teaching: 50 min Exercises: 10 min

Questions

- How can I summarize my data by aggregating, filtering, or ordering query results?
- How can I make sure column names from my queries make sense and aren't too long?

Objectives

- · Apply aggregation to group records in SQL.
- Filter and order results of a query based on aggregate functions.
- Employ aliases to assign new names to items in a query.
- Save a query to make a new table.
- Apply filters to find missing values in SQL.

COUNT and GROUP BY

Aggregation allows us to combine results by grouping records based on value, also it is useful for calculating combined values in groups.

Let's go to the surveys table and find out how many individuals there are. Using the wildcard * simply counts the number of records (rows):

```
SELECT COUNT(*)
FROM surveys;
```

We can also find out how much all of those individuals weigh:

```
SELECT COUNT(*), SUM(weight)
FROM surveys;
```

We can output this value in kilograms (dividing the value to 1000.0), then rounding to 3 decimal places: (Notice the divisor has numbers after the decimal point, which forces the answer to have a decimal fraction)

```
SELECT ROUND(SUM(weight)/1000.00, 3)
FROM surveys;
```

There are many other aggregate functions included in SQL, for example: MAX, MIN, and AVG.

Challenge

Write a query that returns: the total weight, average weight, minimum and maximum weights for all animals caught over the duration of the survey. Can you modify it so that it outputs these values only for weights between 5 and 10?

Solution

```
-- All animals
SELECT SUM(weight), AVG(weight), MIN(weight), MAX(weight)
FROM surveys;

-- Only weights between 5 and 10
SELECT SUM(weight), AVG(weight), MIN(weight), MAX(weight)
FROM surveys
WHERE (weight > 5) AND (weight < 10);
```

Now, let's see how many individuals were counted in each species. We do this using a GROUP BY clause

```
SELECT species_id, COUNT(*)
FROM surveys
GROUP BY species_id;
```

GROUP BY tells SQL what field or fields we want to use to aggregate the data. If we want to group by multiple fields, we give GROUP BY a comma separated list.

```
Challenge
```

Write queries that return:

- 1. How many individuals were counted in each year in total
- 2. How many were counted each year, for each different species
- 3. The average weights of each species in each year

Can you get the answer to both 2 and 3 in a single query?

Solution of 1

```
SELECT year, COUNT(*)
FROM surveys
GROUP BY year;
```

Solution of 2 and 3

```
SELECT year, species_id, COUNT(*), AVG(weight)
FROM surveys
GROUP BY year, species_id;
```

Ordering Aggregated Results

We can order the results of our aggregation by a specific column, including the aggregated column. Let's count the number of individuals of each species captured, ordered by the count:

```
SELECT species_id, COUNT(*)
FROM surveys
GROUP BY species_id
ORDER BY COUNT(species_id);
```

Aliases

As queries get more complex names can get long and unwieldy. To help make things clearer we can use aliases to assign new names to things in the query.

We can use aliases in column names or table names using AS:

```
SELECT MAX(year) AS last_surveyed_year FROM surveys;
```

The AS isn't technically required, so you could do

```
SELECT MAX(year) yr
FROM surveys surv;
```

but using AS is much clearer so it is good style to include it.

The HAVING keyword

In the previous episode, we have seen the keyword <code>WHERE</code>, allowing to filter the results according to some criteria. SQL offers a mechanism to filter the results based on **aggregate functions**, through the <code>HAVING</code> keyword.

For example, we can request to only return information about species with a count higher than 10:

```
SELECT species_id, COUNT(species_id)
FROM surveys
GROUP BY species_id
HAVING COUNT(species_id) > 10;
```

The HAVING keyword works exactly like the WHERE keyword, but uses aggregate functions instead of database fields to filter.

If you use AS in your query to rename a column, HAVING you can use this information to make the query more readable. For example, in the above query, we can call the COUNT(species_id) by another name, like occurrences. This can be written this way:

```
SELECT species_id, COUNT(species_id) AS occurrences
FROM surveys
GROUP BY species_id
HAVING occurrences > 10;
```

Note that in both queries, HAVING comes after GROUP BY. One way to think about this is: the data are retrieved (SELECT), which can be filtered (WHERE), then joined in groups (GROUP BY); finally, we can filter again based on some of these groups (HAVING).

Write a query that returns, from the species table, the number of species in each taxa, only for the taxa with more than 10 species.

Solution

```
SELECT taxa, COUNT(*) AS n
FROM species
GROUP BY taxa
HAVING n > 10;
```

Saving Queries for Future Use

It is not uncommon to repeat the same operation more than once, for example for monitoring or reporting purposes. SQL comes with a very powerful mechanism to do this by creating views. Views are a form of query that is saved in the database, and can be used to look at, filter, and even update information. One way to think of views is as a table, that can read, aggregate, and filter information from several places before showing it to you.

Creating a view from a query requires to add CREATE VIEW viewname AS before the query itself. For example, imagine that my project only covers the data gathered during the summer (May - September) of 2000. That query would look like:

```
SELECT *
FROM surveys
WHERE year = 2000 AND (month > 4 AND month < 10);
```

But we don't want to have to type that every time we want to ask a question about that particular subset of data. Hence, we can benefit from a view:

```
CREATE VIEW summer_2000 AS

SELECT *

FROM surveys

WHERE year = 2000 AND (month > 4 AND month < 10);
```

You can also add a view using *Create View* in the *View* menu and see the results in the *Views* tab, the same way as creating a table from the menu.

Using a view we will be able to access these results with a much shorter notation:

```
SELECT *
FROM summer_2000
WHERE species_id == 'PE';
```

What About NULL?

From the last example, there should only be six records. If you look at the weight column, it's easy to see what the average weight would be. If we use SQL to find the average weight, SQL behaves like we would hope, ignoring the NULL values:

```
SELECT AVG(weight)
FROM summer_2000
WHERE species_id == 'PE';
```

But if we try to be extra clever, and find the average ourselves, we might get tripped up:

```
SELECT SUM(weight), COUNT(*), SUM(weight)/COUNT(*)
FROM summer_2000
WHERE species_id == 'PE';
```

Here the COUNT command includes all six records (even those with NULL values), but the SUM only includes the 4 records with data in the weight field, giving us an incorrect average. However, our strategy *will* work if we modify the COUNT command slightly:

```
SELECT SUM(weight), COUNT(weight), SUM(weight)/COUNT(weight)
FROM summer_2000
WHERE species_id == 'PE';
```

When we count the weight field specifically, SQL ignores the records with data missing in that field. So here is one example where NULLs can be tricky: COUNT(*) and COUNT(field) can return different values.

Another case is when we use a "negative" query. Let's count all the non-female animals:

```
SELECT COUNT(*)
FROM summer_2000
WHERE sex != 'F';
```

Now let's count all the non-male animals:

```
SELECT COUNT(*)
FROM summer_2000
WHERE sex != 'M';
```

But if we compare those two numbers with the total:

```
SELECT COUNT(*)
FROM summer_2000;
```

We'll see that they don't add up to the total! That's because SQL doesn't automatically include NULL values in a negative conditional statement. So if we are quering "not x", then SQL divides our data into three categories: 'x', 'not NULL, not x' and NULL; then, returns the 'not NULL, not x' group. Sometimes this may be what we want - but sometimes we may want the missing values included as well! In that case, we'd need to change our query to:

```
SELECT COUNT(*)
FROM summer_2000
WHERE sex != 'M' OR sex IS NULL;
```

Rey Points

- Use the GROUP BY keyword to aggregate data.
- Functions like MIN, MAX, AVERAGE, SUM, COUNT, etc. operate on aggregated data.
- Aliases can help shorten long queries. To write clear and readible queries, use the AS keyword when creating aliases.
- Use the HAVING keyword to filter on aggregate properties.
- Use a VIEW to access the result of a query as though it was a new table.

Joins



Teaching: 15 min **Exercises:** 10 min

Questions

How do I bring data together from separate tables?

Objectives

- Employ joins to combine data from two tables.
- · Apply functions to manipulate individual values.
- Employ aliases to assign new names to tables and columns in a query.

Joins

To combine data from two tables we use the SQL JOIN command, which comes after the FROM command.

The JOIN command on its own will result in a cross product, where each row in the first table is paired with each row in the second table. Usually this is not what is desired when combining two tables with data that is related in some way.

For that, we need to tell the computer which columns provide the link between the two tables using the word ON. What we want is to join the data with the same species id.

```
SELECT *
FROM surveys
JOIN species
ON surveys.species_id = species.species_id;
```

ON is like WHERE, it filters things out according to a test condition. We use the table.colname format to tell the manager what column in which table we are referring to.

The output of the JOIN command will have columns from the first table plus the columns from the second table. For the above command, the output will be a table that has the following column names:

 $record_idmonth day year\ plot_idspecies_idsex hind foot_length weight species_idgenus \qquad species\ tax a discontinuous and a species between the property of t$

96 8 20 199712 **DM** M 36 41 **DM** DipodomysmerriamiRodent

Alternatively, we can use the word USING, as a short-hand. USING only works on columns which share the same name. In this case we are telling the manager that we want to combine surveys with species and that the common column is species_id.

```
SELECT *
FROM surveys
JOIN species
USING (species_id);
```

The output will only have one **species_id** column

record_idmonthdayyear plot_idspecies_idsexhindfoot_lengthweightgenus species taxa

• • •

96 8 20 199712 DM M 36 41 DipodomysmerriamiRodent

record_idmonthdayyear plot_idspecies_idsexhindfoot_lengthweightgenus species taxa

. . .

We often won't want all of the fields from both tables, so anywhere we would have used a field name in a non-join query, we can use table.colname.

For example, what if we wanted information on when individuals of each species were captured, but instead of their species ID we wanted their actual species names.

```
SELECT surveys.year, surveys.month, surveys.day, species.genus, species.species
FROM surveys
JOIN species
ON surveys.species_id = species.species_id;
```

year monthday genus species

...

19777 16 Neotoma albigula19777 16 Dipodomysmerriami

. . .

Many databases, including SQLite, also support a join through the WHERE clause of a query. For example, you may see the query above written without an explicit JOIN.

```
SELECT surveys.year, surveys.month, surveys.day, species.genus, species.species
FROM surveys, species
WHERE surveys.species_id = species.species_id;
```

For the remainder of this lesson, we'll stick with the explicit use of the JOIN keyword for joining tables in SQL.

Challenge:

 Write a query that returns the genus, the species name, and the weight of every individual captured at the site

Solution

```
SELECT species.genus, species.species, surveys.weight
FROM surveys
JOIN species
ON surveys.species_id = species.species_id;
```

Different join types

We can count the number of records returned by our original join query.

```
SELECT COUNT(*)
FROM surveys
JOIN species
USING (species_id);
```

Notice that this number is smaller than the number of records present in the survey data.

```
SELECT COUNT(*) FROM surveys;
```

This is because, by default, SQL only returns records where the joining value is present in the joined columns of both tables (i.e. it takes the *intersection* of the two join columns). This joining behaviour is known as an INNER JOIN . In fact the JOIN command is simply shorthand for INNER JOIN and the two terms can be used interchangably as they will produce the same result.

We can also tell the computer that we wish to keep all the records in the first table by using the command LEFT OUTER JOIN, or LEFT JOIN for short.

Challenge:

 Re-write the original query to keep all the entries present in the surveys table. How many records are returned by this query?

Solution

```
SELECT * FROM surveys
LEFT JOIN species
USING (species_id);
```

Challenge:

Count the number of records in the surveys table that have a NULL value in the species_id column.

Solution

```
SELECT COUNT(*)
FROM surveys
WHERE species_id IS NULL;
```

Remember: In SQL a NULL value in one table can never be joined to a NULL value in a second table because NULL is not equal to anything, not even itself.

Combining joins with sorting and aggregation

Joins can be combined with sorting, filtering, and aggregation. So, if we wanted average mass of the individuals on each different type of treatment, we could do something like

```
SELECT plots.plot_type, AVG(surveys.weight)
FROM surveys
JOIN plots
ON surveys.plot_id = plots.plot_id
GROUP BY plots.plot_type;
```

Challenge:

• Write a query that returns the number of animals caught of each genus in each plot. Order the results by plot number (ascending) and by descending number of individuals in each plot.

Solution

```
SELECT surveys.plot_id, species.genus, COUNT(*) AS number_indiv
FROM surveys
JOIN species
ON surveys.species_id = species.species_id
GROUP BY species.genus, surveys.plot_id
ORDER BY surveys.plot_id ASC, number_indiv DESC;
```

Challenge:

• Write a query that finds the average weight of each rodent species (i.e., only include species with Rodent in the taxa field).

Solution

```
SELECT surveys.species_id, AVG(surveys.weight)
FROM surveys
JOIN species
ON surveys.species_id = species.species_id
WHERE species.taxa = 'Rodent'
GROUP BY surveys.species_id;
```

Functions IFNULL and NULLIF and more

SQL includes numerous functions for manipulating data. You've already seen some of these being used for aggregation (SUM and COUNT) but there are functions that operate on individual values as well. Probably the most important of these are IFNULL and NULLIF. IFNULL allows us to specify a value to use in place of NULL.

We can represent unknown sexes with 'U' instead of NULL:

```
SELECT species_id, sex, IFNULL(sex, 'U')
FROM surveys;
```

The lone "sex" column is only included in the query above to illustrate where IFNULL has changed values; this isn't a usage requirement.

Challenge:

• Write a query that returns 30 instead of NULL for values in the hindfoot_length column.

Solution

```
SELECT hindfoot_length, IFNULL(hindfoot_length, 30)
FROM surveys;
```

Challenge:

• Write a query that calculates the average hind-foot length of each species, assuming that unknown lengths are 30 (as above).

Solution

```
SELECT species_id, AVG(IFNULL(hindfoot_length,30))
FROM surveys
GROUP BY species_id;
```

IFNULL can be particularly useful in JOIN. When joining the species and surveys tables earlier, some results were excluded because the species_id was NULL in the surveys table. We can use IFNULL to include them again, re-writing the NULL to a valid joining value:

```
SELECT surveys.year, surveys.month, surveys.day, species.genus, species.species
FROM surveys
JOIN species
ON IFNULL(surveys.species_id, 'AB') = species.species_id;
```

Challenge:

• Write a query that returns the number of animals caught of each genus in each plot, using IFNULL to assume that unknown species are all of the genus "Rodent".

Solution

```
SELECT plot_id, IFNULL(genus, 'Rodent') AS genus2, COUNT(*)
FROM surveys
LEFT JOIN species
ON surveys.species_id=species.species_id
GROUP BY plot_id, genus2;
```

The inverse of IFNULL is NULLIF. This returns NULL if the first argument is equal to the second argument. If the two are not equal, the first argument is returned. This is useful for "nulling out" specific values.

We can "null out" plot 7:

```
SELECT species_id, plot_id, NULLIF(plot_id, 7)
FROM surveys;
```

Some more functions which are common to SQL databases are listed in the table below:

Function	Description
ABS(n)	Returns the absolute (positive) value of the numeric expression n
LENGTH(s)	Returns the length of the string expression s
LOWER(s)	Returns the string expression s converted to lowercase
NULLIF(x, y)	Returns NULL if x is equal to y , otherwise returns x
ROUND(n) or	Returns the numeric expression n rounded to x digits after the decimal point (0 by
ROUND(n, x)	default)
TRIM(s)	Returns the string expression s without leading and trailing whitespace characters
UPPER(s)	Returns the string expression s converted to uppercase

Finally, some useful functions which are particular to SQLite are listed in the table below:

Function	Description
IFNULL(x, y)	Returns <i>x</i> if it is non-NULL, otherwise returns <i>y</i>
RANDOM()	Returns a random integer between -9223372036854775808 and +9223372036854775807.
REPLACE(s, f, r)	Returns the string expression s in which every occurrence of f has been replaced with r
SUBSTR(s, x, y) or	Returns the portion of the string expression s starting at the character position x (leftmost
SUBSTR(s, x)	position is 1), y characters long (or to the end of s if y is omitted)

Challenge:

Write a query that returns genus names (no repeats), sorted from longest genus name down to shortest.

Solution

SELECT DISTINCT genus FROM species ORDER BY LENGTH(genus) DESC;

As we saw before, aliases make things clearer, and are especially useful when joining tables.

```
SELECT surv.year AS yr, surv.month AS mo, surv.day AS day, sp.genus AS gen, sp.species A
S sp
FROM surveys AS surv
JOIN species AS sp
ON surv.species_id = sp.species_id;
```

To practice we have some optional challenges for you.

Challenge (optional):

SQL queries help us *ask* specific *questions* which we want to answer about our data. The real skill with SQL is to know how to translate our scientific questions into a sensible SQL query (and subsequently visualize and interpret our results).

Have a look at the following questions; these questions are written in plain English. Can you translate them to *SQL queries* and give a suitable answer?

- 1. How many plots from each type are there?
- 2. How many specimens are of each sex are there for each year?
- 3. How many specimens of each species were captured in each type of plot?
- 4. What is the average weight of each taxa?
- 5. What are the minimum, maximum and average weight for each species of Rodent?
- 6. What is the average hindfoot length for male and female rodent of each species? Is there a Male / Female difference?
- 7. What is the average weight of each rodent species over the course of the years? Is there any noticeable trend for any of the species?

Proposed solutions:

1. Solution:

```
SELECT plot_type, COUNT(*) AS num_plots
FROM plots
GROUP BY plot_type;
```

1. Solution:

```
SELECT year, sex, COUNT(*) AS num_animal FROM surveys
WHERE sex IS NOT NULL
GROUP BY sex, year;
```

1. Solution:

```
SELECT species_id, plot_type, COUNT(*)
FROM surveys
JOIN plots USING(plot_id)
WHERE species_id IS NOT NULL
GROUP BY species_id, plot_type;
```

1. Solution:

```
SELECT taxa, AVG(weight)
FROM surveys
JOIN species ON species_id = surveys.species_id
GROUP BY taxa;
```

1. Solution:

```
SELECT surveys.species_id, MIN(weight), MAX(weight), AVG(weight) FROM surveys
JOIN species ON surveys.species_id = species.species_id
WHERE taxa = 'Rodent'
GROUP BY surveys.species_id;
```

1. Solution:

```
SELECT surveys.species_id, sex, AVG(hindfoot_length)
FROM surveys JOIN species ON surveys.species_id = species.species_id
WHERE (taxa = 'Rodent') AND (sex IS NOT NULL)
GROUP BY surveys.species_id, sex;
```

1. Solution:

```
SELECT surveys.species_id, year, AVG(weight) as mean_weight FROM surveys

JOIN species ON surveys.species_id = species.species_id

WHERE taxa = 'Rodent' GROUP BY surveys.species_id, year;
```

• Key Points

- Use the JOIN command to combine data from two tables—the ON or USING keywords specify which
 columns link the tables.
- Regular JOIN returns only matching rows. Other join commands provide different behavior, e.g.,
 LEFT JOIN retains all rows of the table on the left side of the command.
- IFNULL allows you to specify a value to use in place of NULL, which can help in joins
- NULLIF can be used to replace certain values with NULL in results
- Many other functions like IFNULL and NULLIF can operate on individual values.

Copyright © 2018–2018 The Carpentries (https://carpentries.org/) Copyright © 2016–2018 Data Carpentry (http://datacarpentry.org)

Edit on GitHub (https://github.com/datacarpentry/sql-ecology-lesson/edit/gh-pages/aio.md) / Contributing (https://github.com/datacarpentry/sql-ecology-lesson/blob/gh-pages/CONTRIBUTING.md) / Source (https://github.com/datacarpentry/sql-ecology-lesson/) / Cite (https://github.com/datacarpentry/sql-ecology-lesson/blob/gh-pages/CITATION) / Contact (mailto:team@carpentries.org)

Using The Carpentries style (https://github.com/carpentries/styles/) version 9.5.2 (https://github.com/carpentries/styles/releases/tag/v9.5.2).